
CSC 660: Advanced OS

Netfilter

Topics

1. What is a firewall?
2. Packet Filtering with iptables
3. Netfilter Architecture
4. Packet Data Structures
5. Netfilter Data Structures
6. Writing a Netfilter extension

What is a Firewall?

- A software or hardware component that restricts network communication between two computers or networks
- In buildings, a firewall is a fireproof wall that restricts the spread of a fire
 - Network firewall prevents threats from spreading from one network to another

What is a Firewall? (2)

- A mechanism to enforce security policy
 - Choke point that traffic has to flow through
 - ACLs on a host/network level
- Policy Decisions:
 - What traffic should be allowed into network?
 - Integrity: protect integrity of internal systems
 - Availability: protection from DOS attacks
 - What traffic should be allowed out of network?
 - Confidentiality: protection from data leakage

Packet Filtering

- Forward or drop packets based on TCP/IP header information, most often:
 - IP source and destination addresses
 - Protocol (ICMP, TCP, or UDP)
 - TCP/UDP source and destination ports
 - TCP Flags, especially SYN and ACK
 - ICMP message type
- Dual-homed hosts also make decisions based on:
 - Network interface the packet arrived on
 - Network interface the packet will depart on

iptables

- Linux packet filtering system
 - iptables is user command to configure
 - netfilter is internal kernel architecture
- Features
 - Packet filtering
 - Connection tracking
 - Network Address Translation

Tables

Each table is a named array of rules.

Within a table, rules are organized into chains.

Tables:

- **Filter**

- Packet filtering (no alterations), default table.
- Hooks: LOCAL_IN, LOCAL_OUT, FORWARD

- **NAT**

- Network address translation.
- Hooks: LOCAL_OUT, PREROUTING, POSTROUTING

- **Mangle**

- Flexible packet alterations.
- Hooks: LOCAL_OUT, PREROUTING

iptables

iptables [-t table] cmd [matches] [target]

Commands:

- A chain rule-spec: Append rule to chain.
- D chain rule-spec: Delete a rule from chain
- L chain: List all rules in chain.
- F chain: Flush all rules from chain.
- P chain target: Set default policy for chain.
- N chain: Create a new chain.
- X chain: Remove a user-defined chain.

iptables Matches

- p protocol:** Specify protocol to match.
tcp, udp, icmp, etc.
- s address/mask:** Source IP address to match.
- d address/mask:** Dest IP address to match.
- sport:** Source port (TCP/UDP) to match.
- dport:** Dest port (TCP/UDP) to match.

iptables Extended Matches

-m match: Specify match module to use.

Example: limit

Only accept 3 ICMP packets per hour.

`-m limit --limit 3/hour -p icmp -j REJECT`

Example: state

Useful stateful packet filtering.

`-m state --state NEW:` match only new conns

`-m state --state ESTABLISHED:` match only established connections.

iptables Targets

-j ACCEPT

Accept packet.

-j DROP

Drop packet w/o reply.

-j REJECT

Drop packet with reply.


-j RETURN

Return from this chain to calling chain.

-j LOG

Log packet; chain processing continues.

Chain Targets



Creating a Packet Filter

1. Create a security policy for a service.
ex: allow only outgoing telnet service
2. Specify security policy in terms of which types of packets are allowed/forbidden
3. Write packet filter in terms of vendor's filtering language

Example: outgoing telnet

- TCP-based service
- Outbound packets
 - Destination port is 23
 - Source port is random port >1023
 - Outgoing connection established by first packet with no ACK flag set
 - Following packets will have ACK flag set
- Incoming packets
 - Source port is 23, as server runs on port 23
 - Destination port is high port used for outbound packets
 - All incoming packets will have ACK flag set

Example: outgoing telnet

- First rule allows outgoing telnet packets
- Second rule allows response packets back in
- Third rule denies all else, following Principle of Fail-Safe Defaults

Dir	Src	Dest	Proto	S.Port	D.Port	ACK?	Action
Out	Int	Any	TCP	>1023	23	Either	Accept
In	Any	Int	TCP	23	>1023	Yes	Accept
Either	Any	Any	Any	Any	Any	Either	Deny

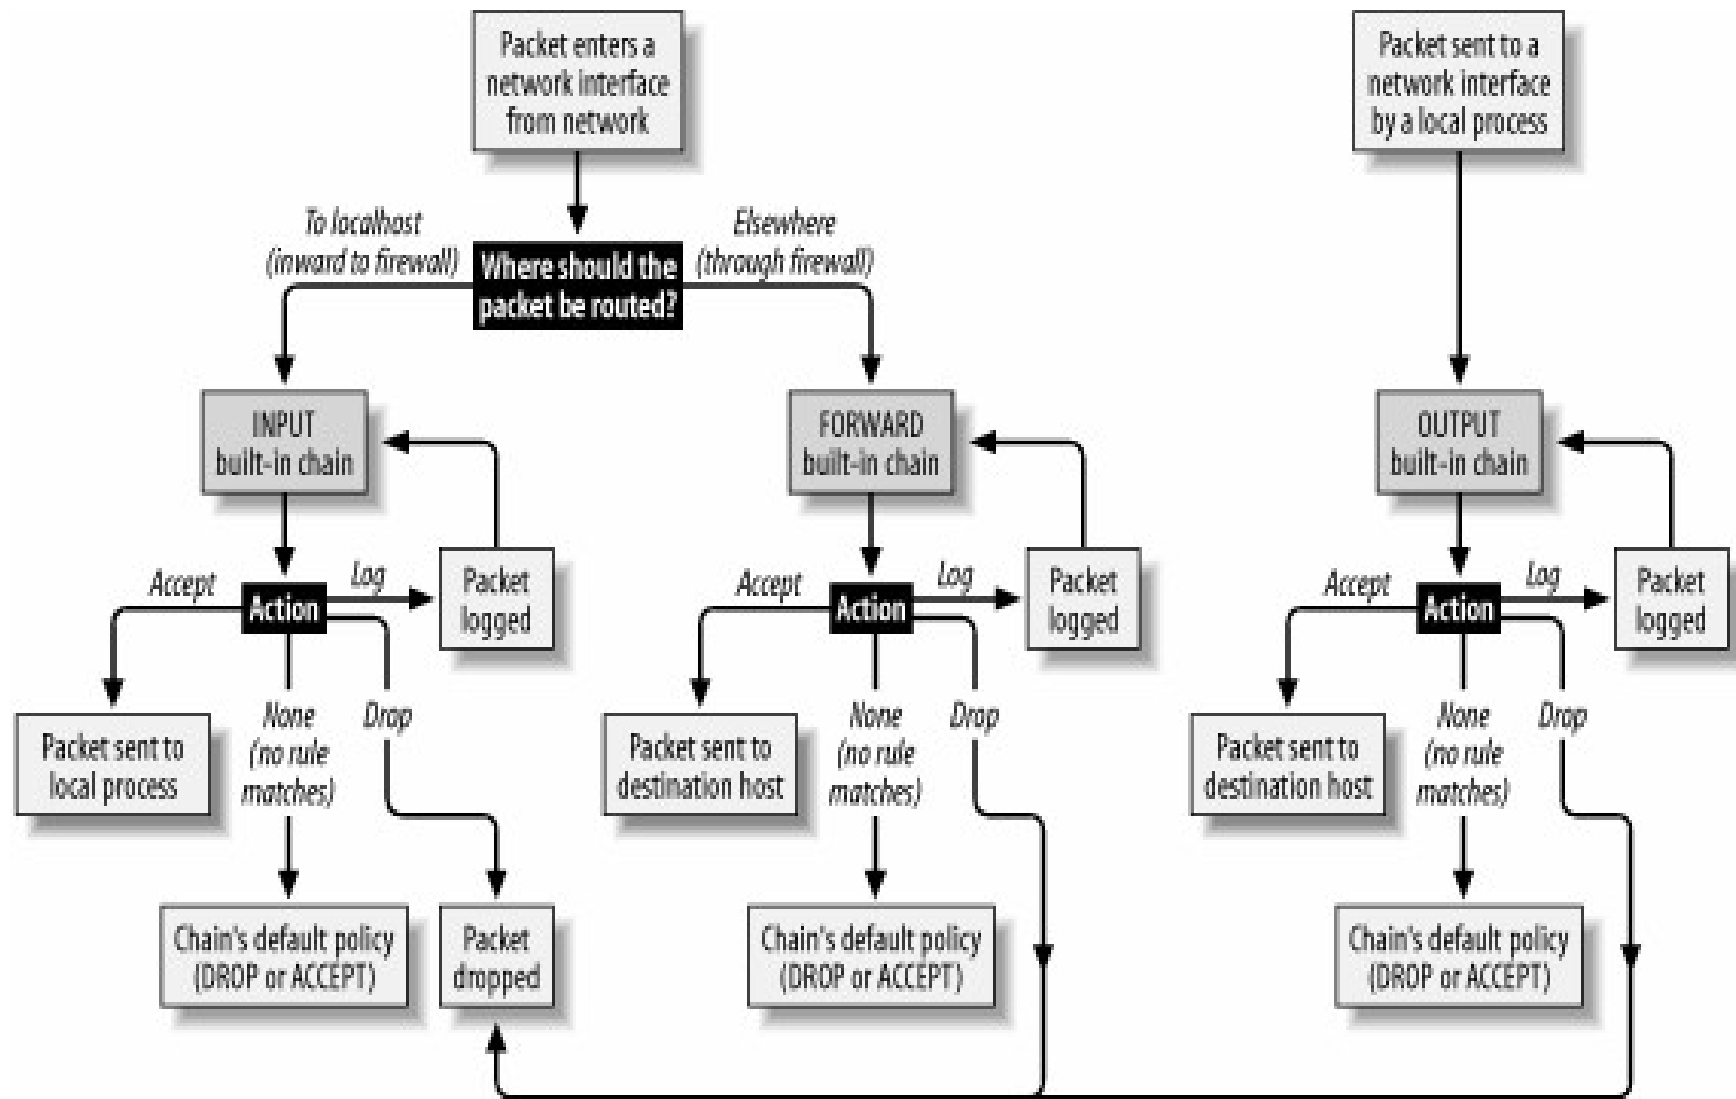
Implementing the Filter with iptables

```
# iptables -A INPUT -m state --state NEW -m  
tcp -p tcp --dport 23 -j ACCEPT
```

```
# iptables -A INPUT -m state --state  
ESTABLISHED,RELATED -m tcp -d tcp  
--sport 23 -j ACCEPT
```

```
# iptables -A INPUT -j REJECT
```


Packet Filtering Hooks



Netfilter Hooks

LOCAL_OUT

Hook for outgoing packets that are created locally.

LOCAL_IN

In `ip_local_deliver()` for incoming pkts destined for localhost.

PRE_ROUTING

Hook for incoming packets in `ip_rcv()` before routing.

FORWARD

In `ip_forward()` for incoming packets destined for another host.

POST_ROUTING

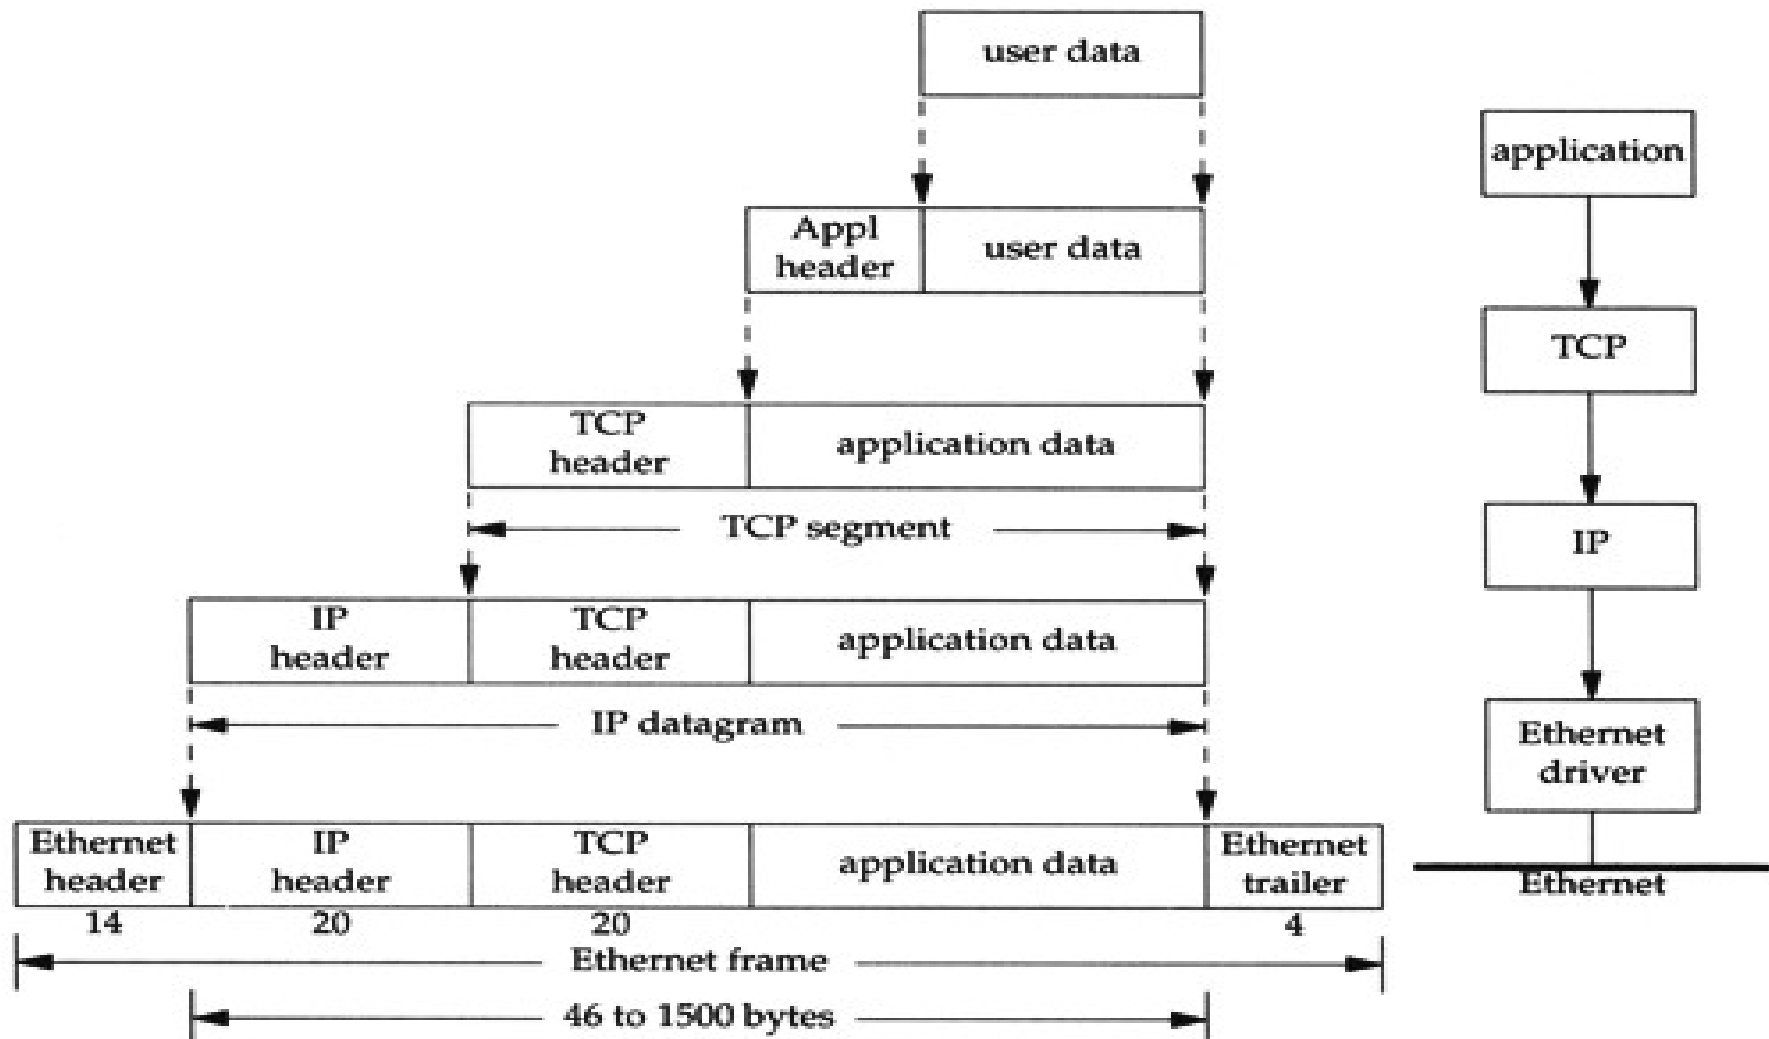
Hook in `ip_finish_output()` for all outgoing packets.

TCP/IP Layering

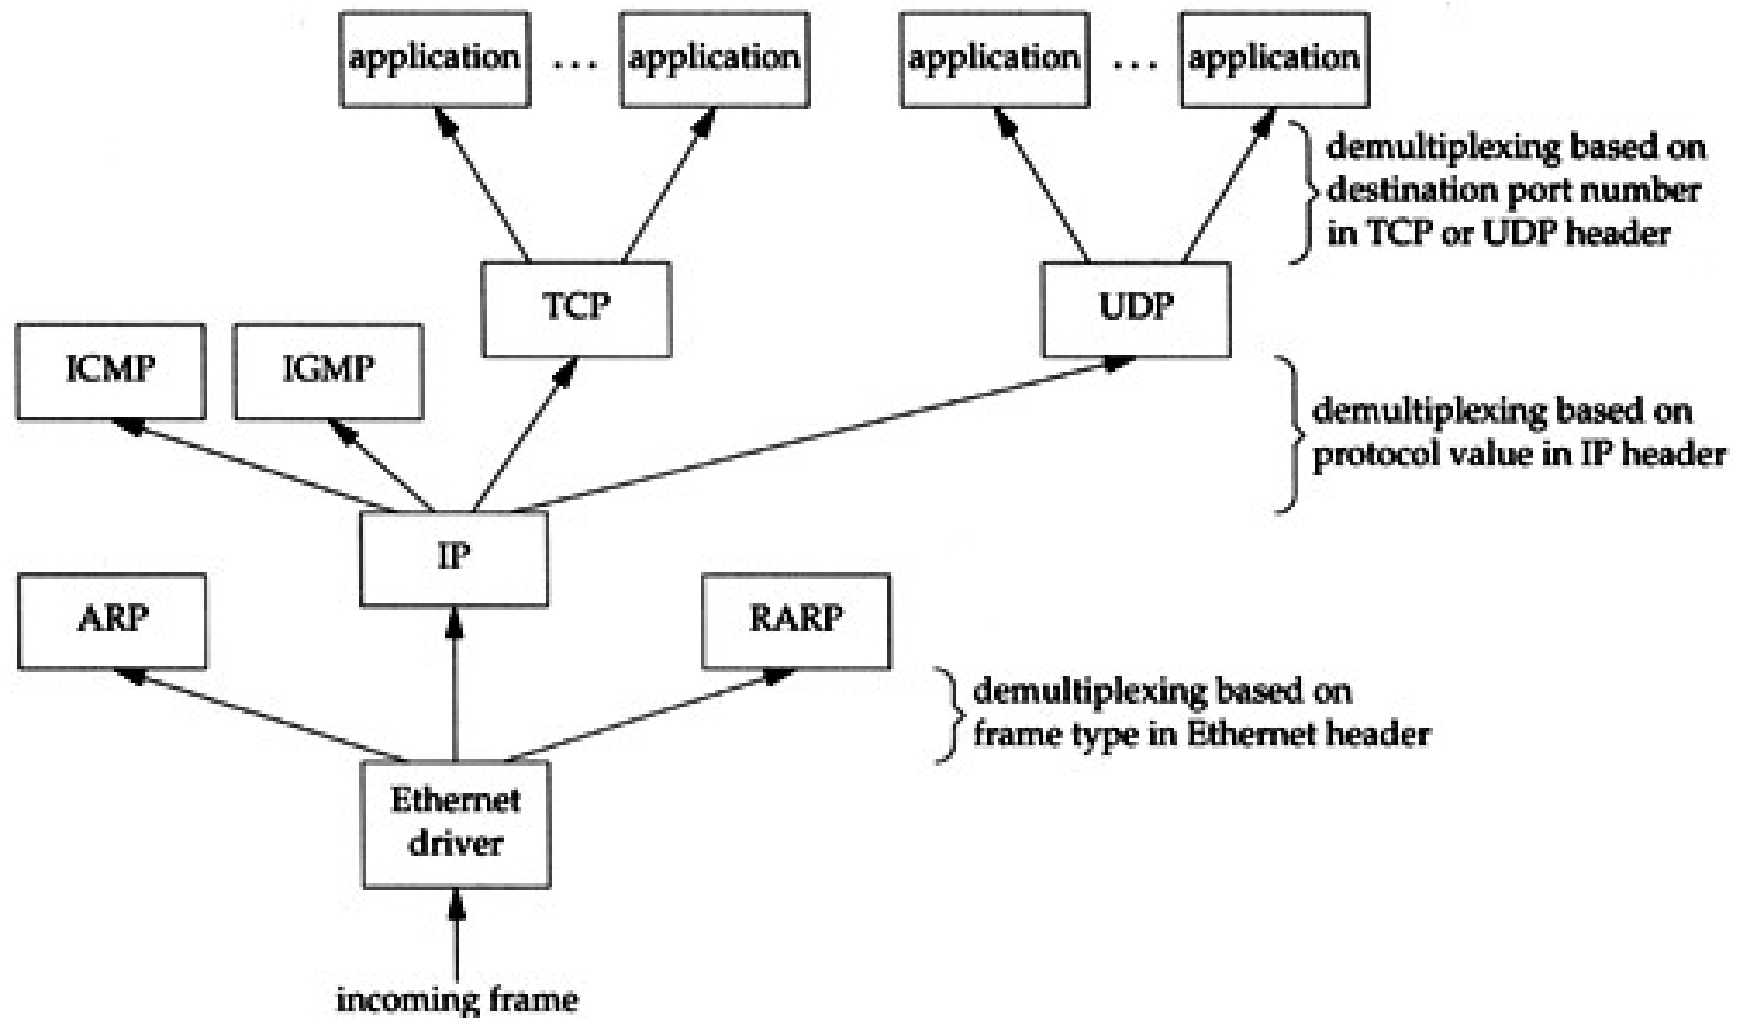
Application
Transport
Network
Data Link

- HTTP, FTP, telnet
- TCP, UDP
- IP, ICMP, IGMP
- Ethernet, PPP, 802.11

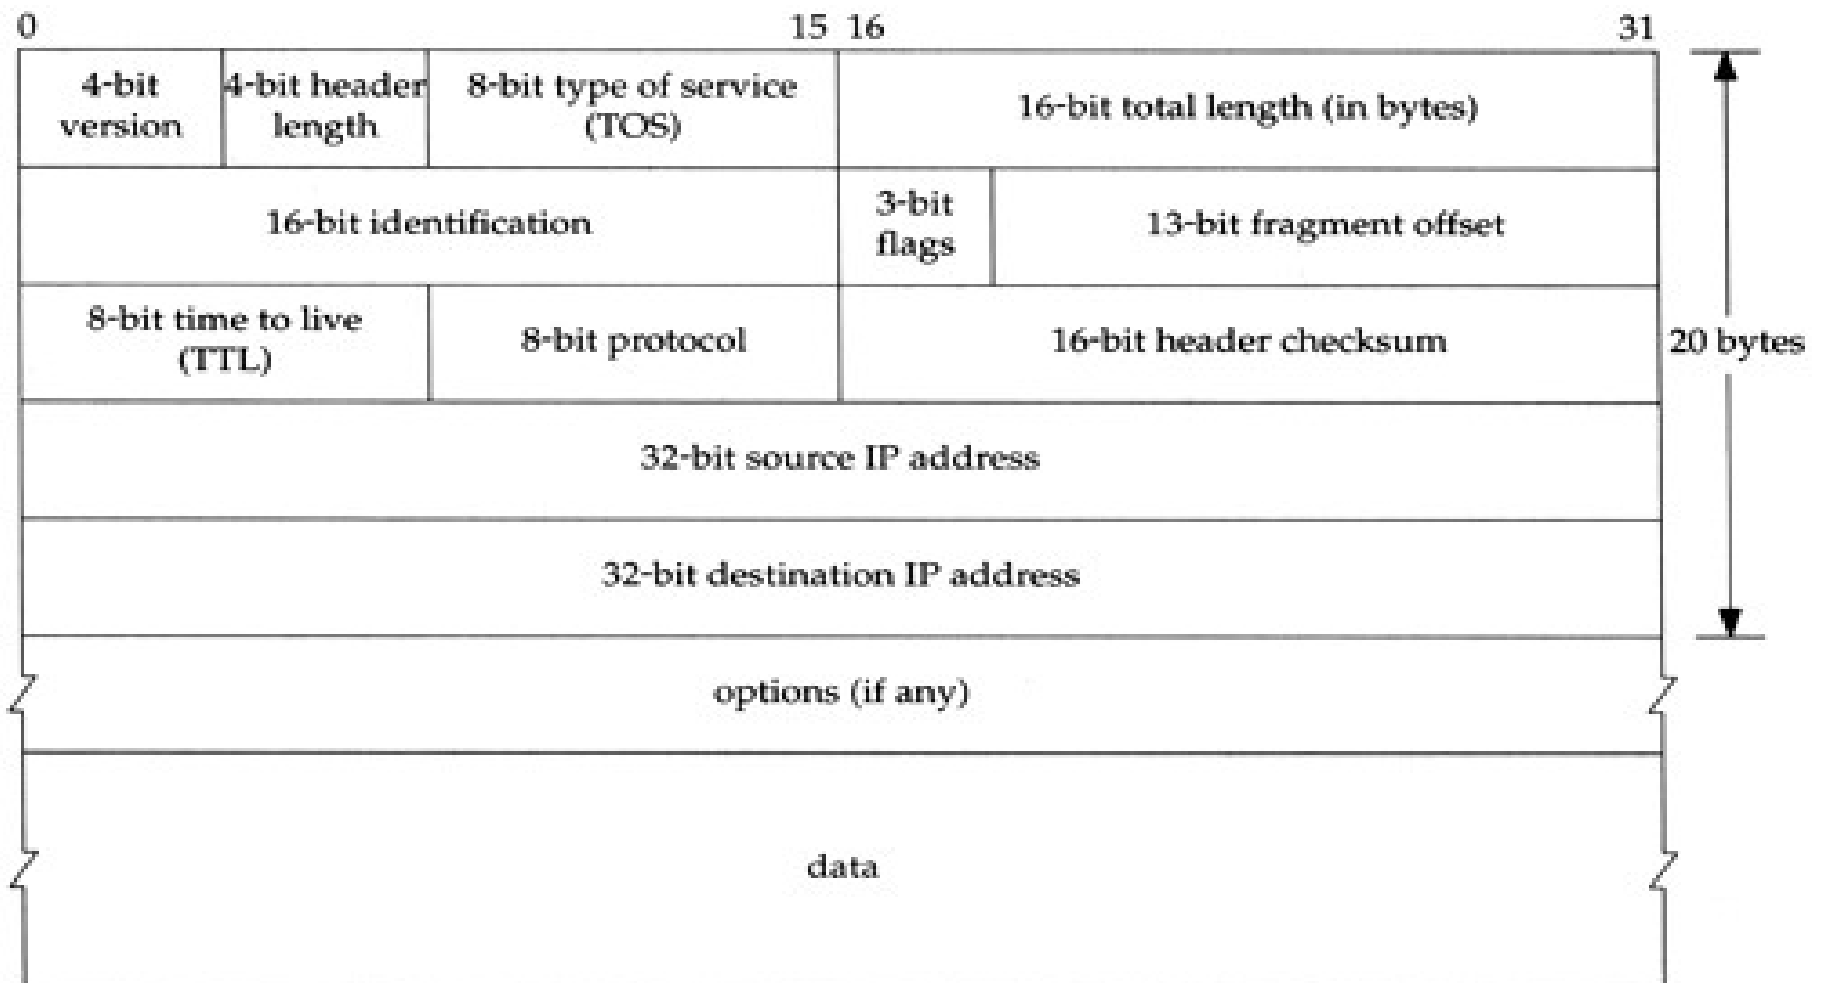
Packet Encapsulation



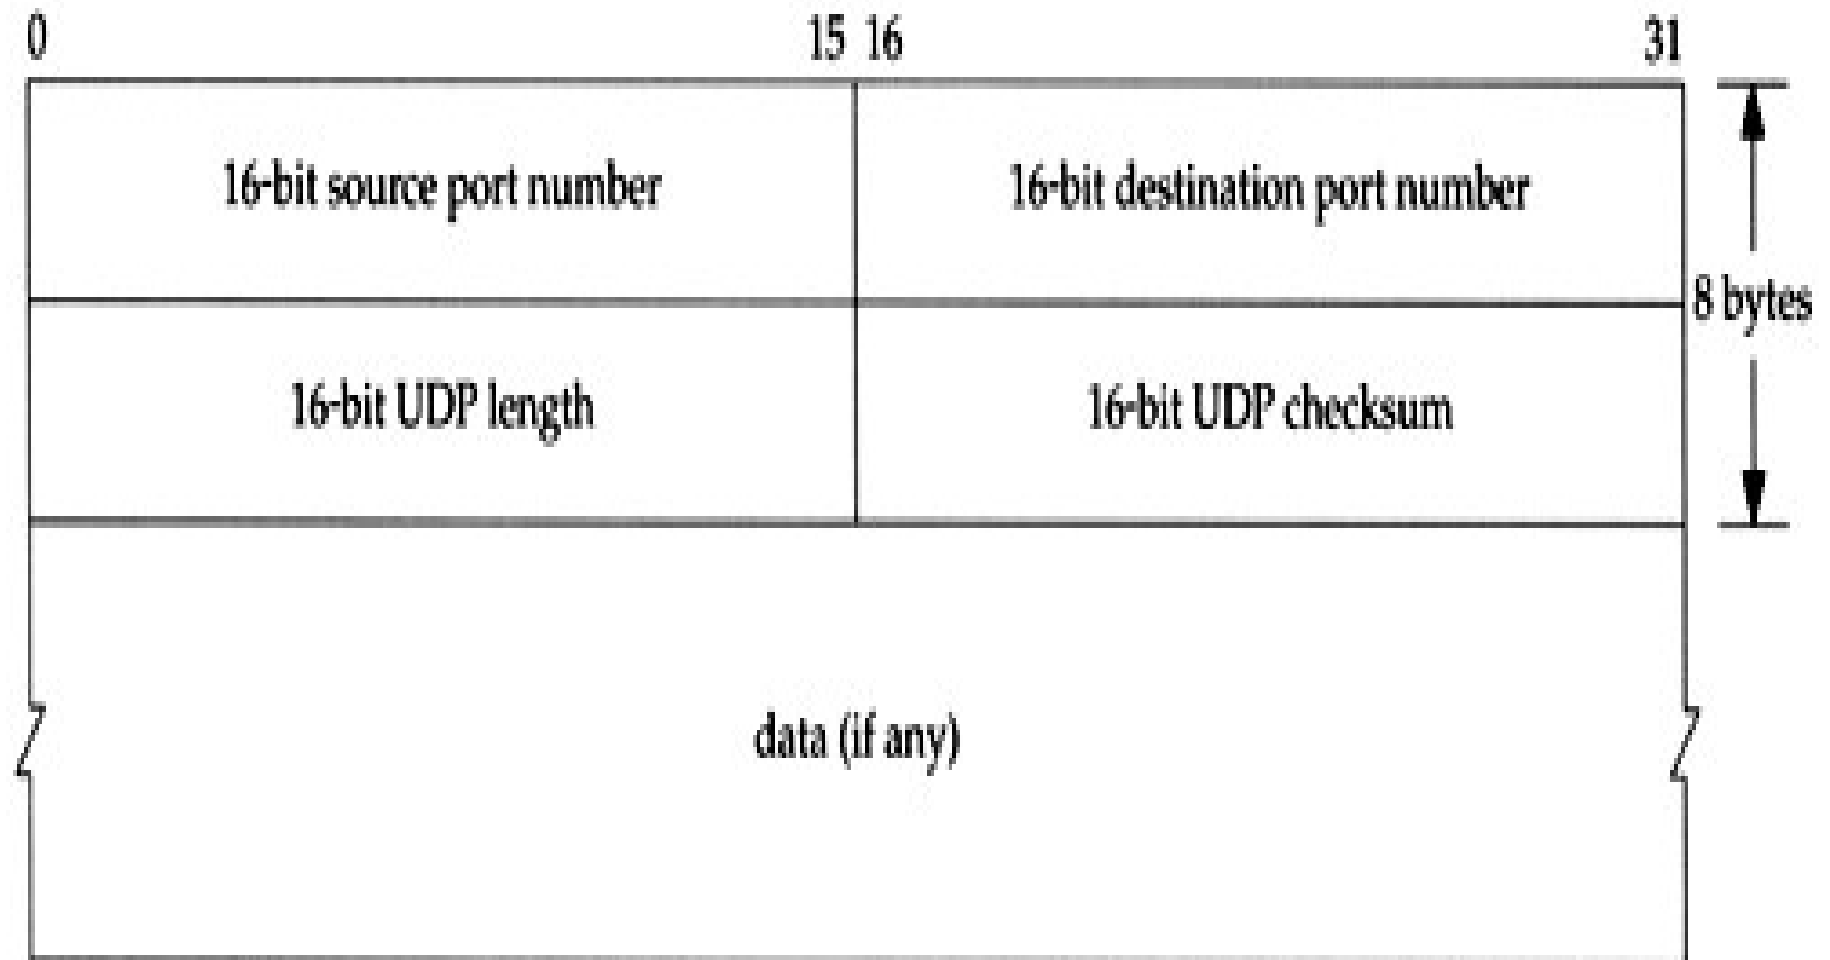
Packet De-multiplexing



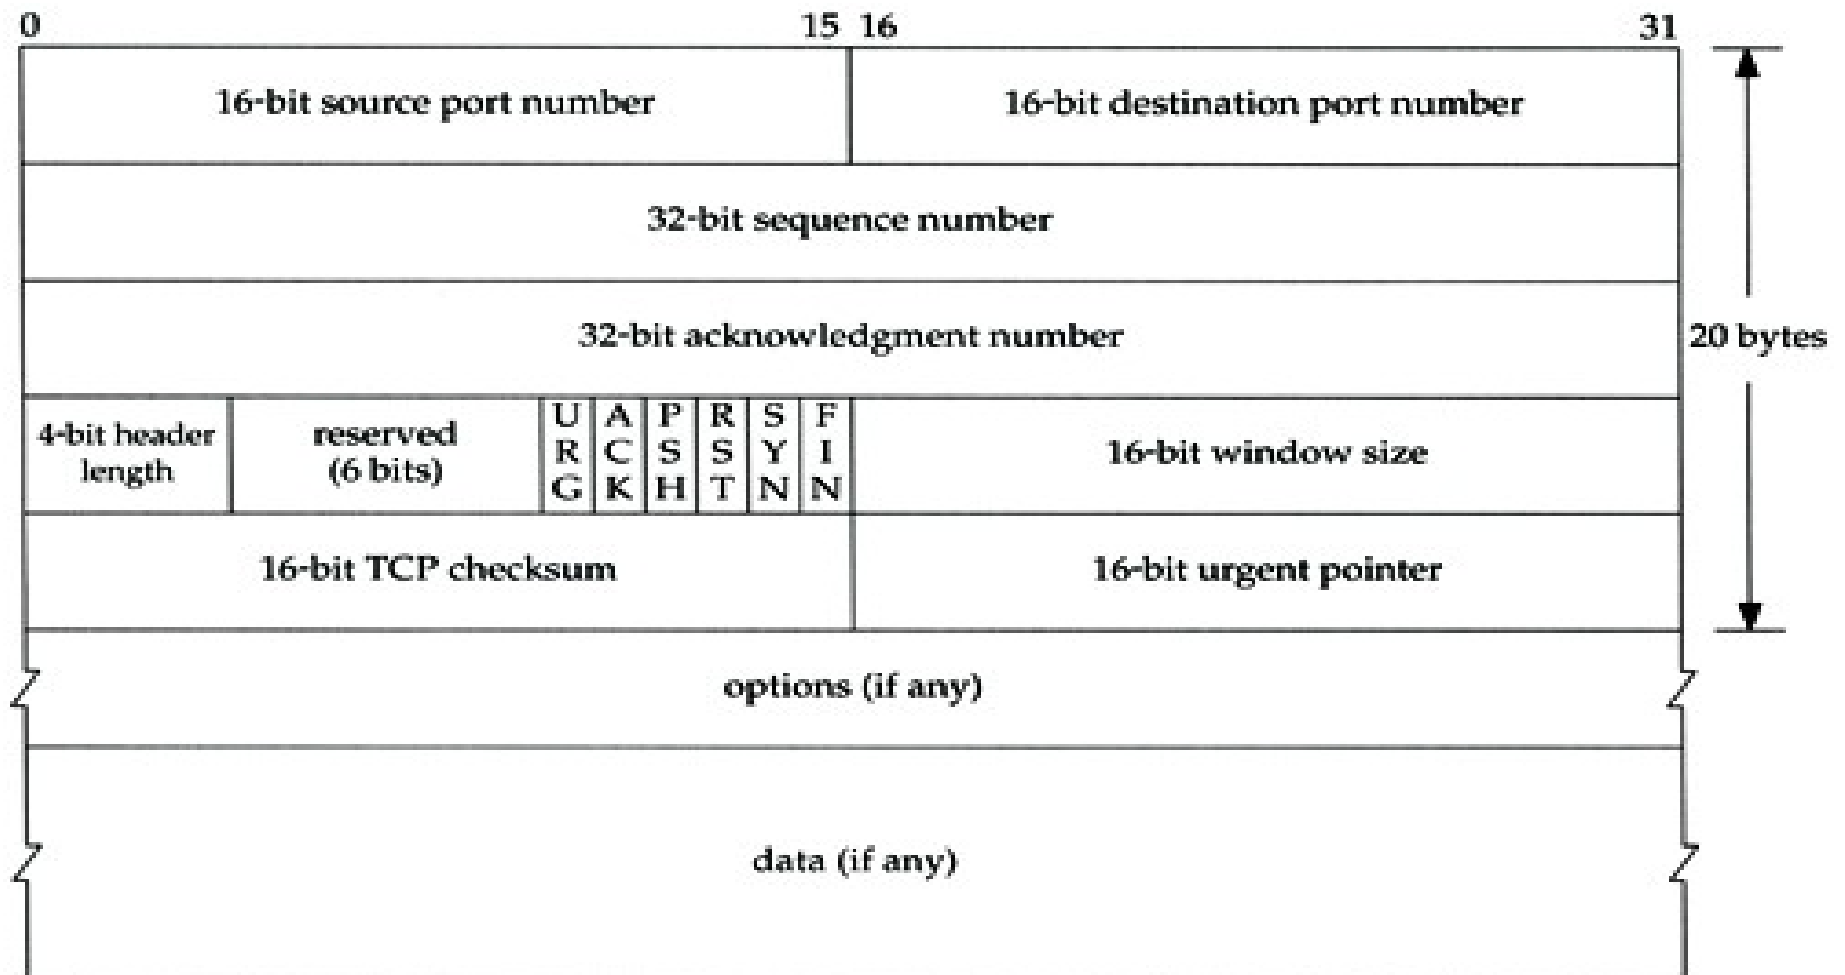
IP Header



UDP Header



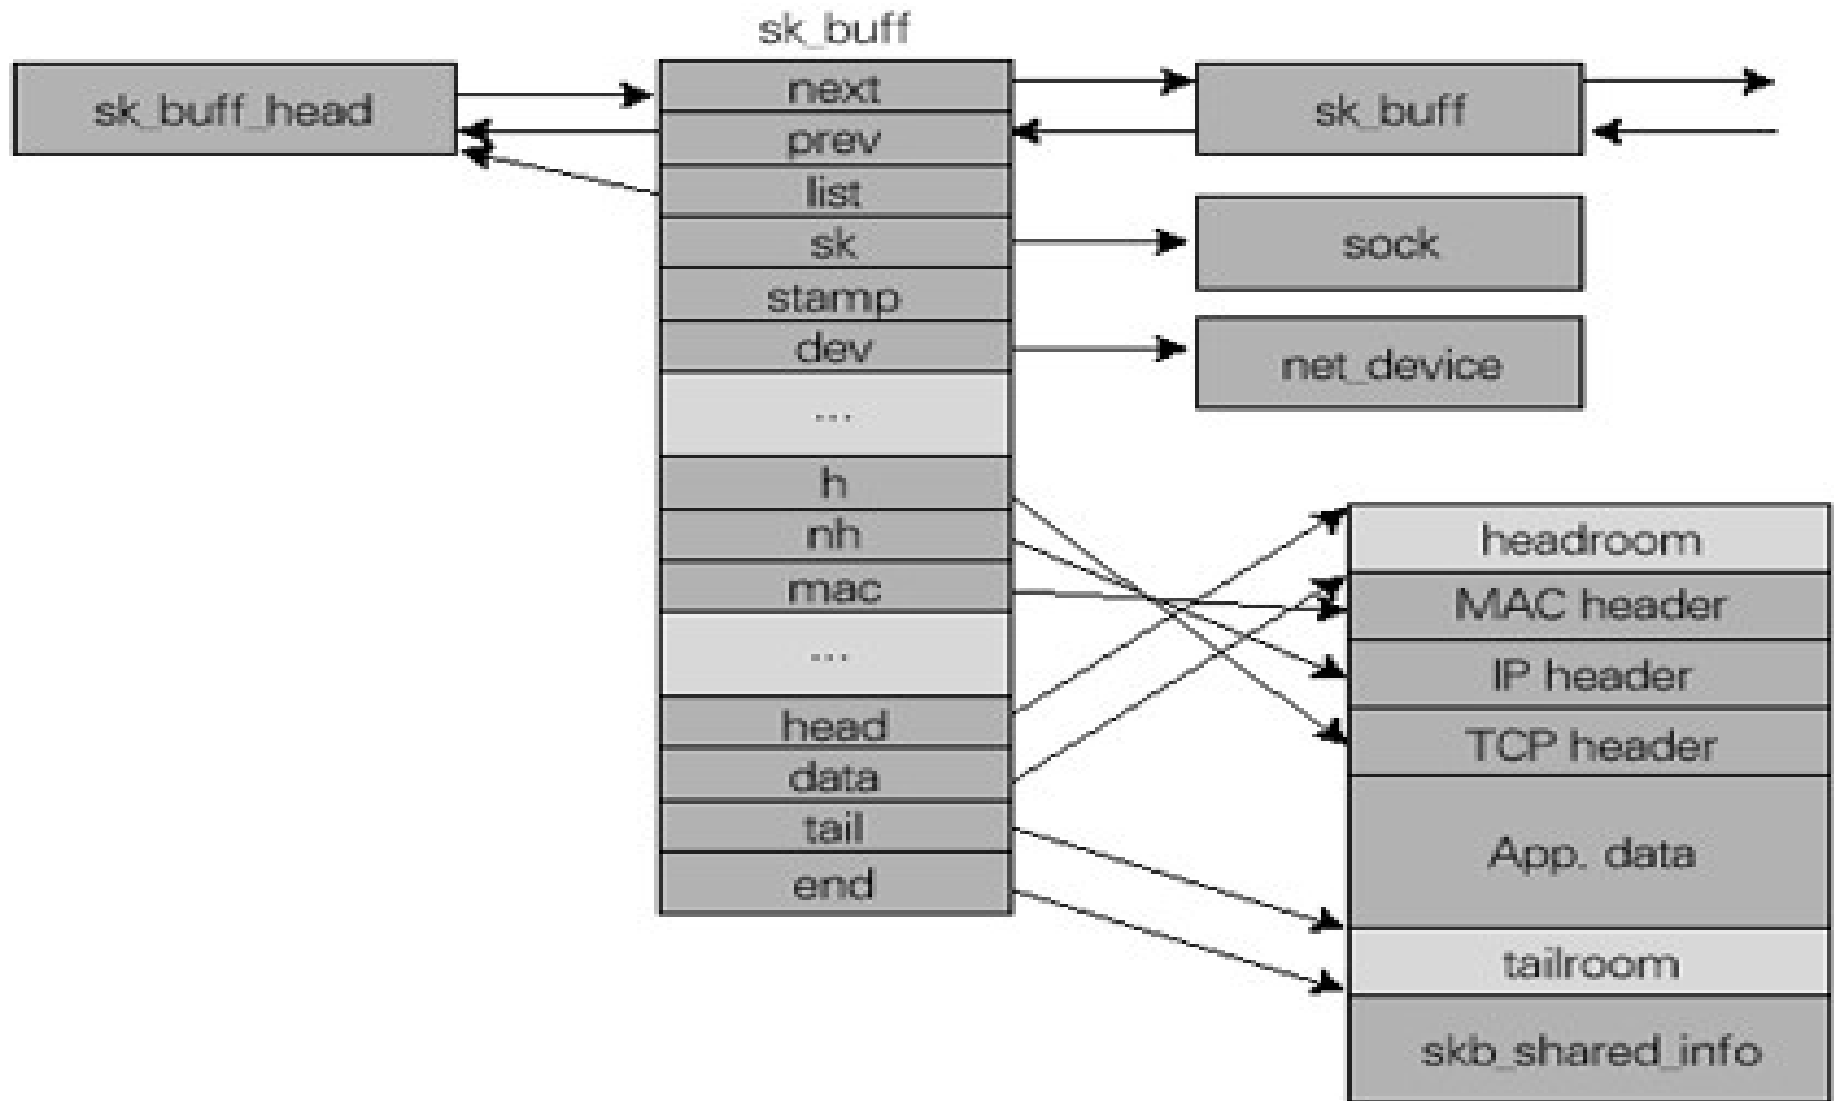
TCP Header



sk_buff

- Kernel buffer that stores packets.
 - Contains headers for all network layers.
- Creation
 - Application sends data to socket.
 - Packet arrives at network interface.
- Copying
 - Copied from user/kernel space.
 - Copied from kernel space to NIC.
 - Send: appends headers via `skb_reserve()`.
 - Receive: moves ptr from header to header.

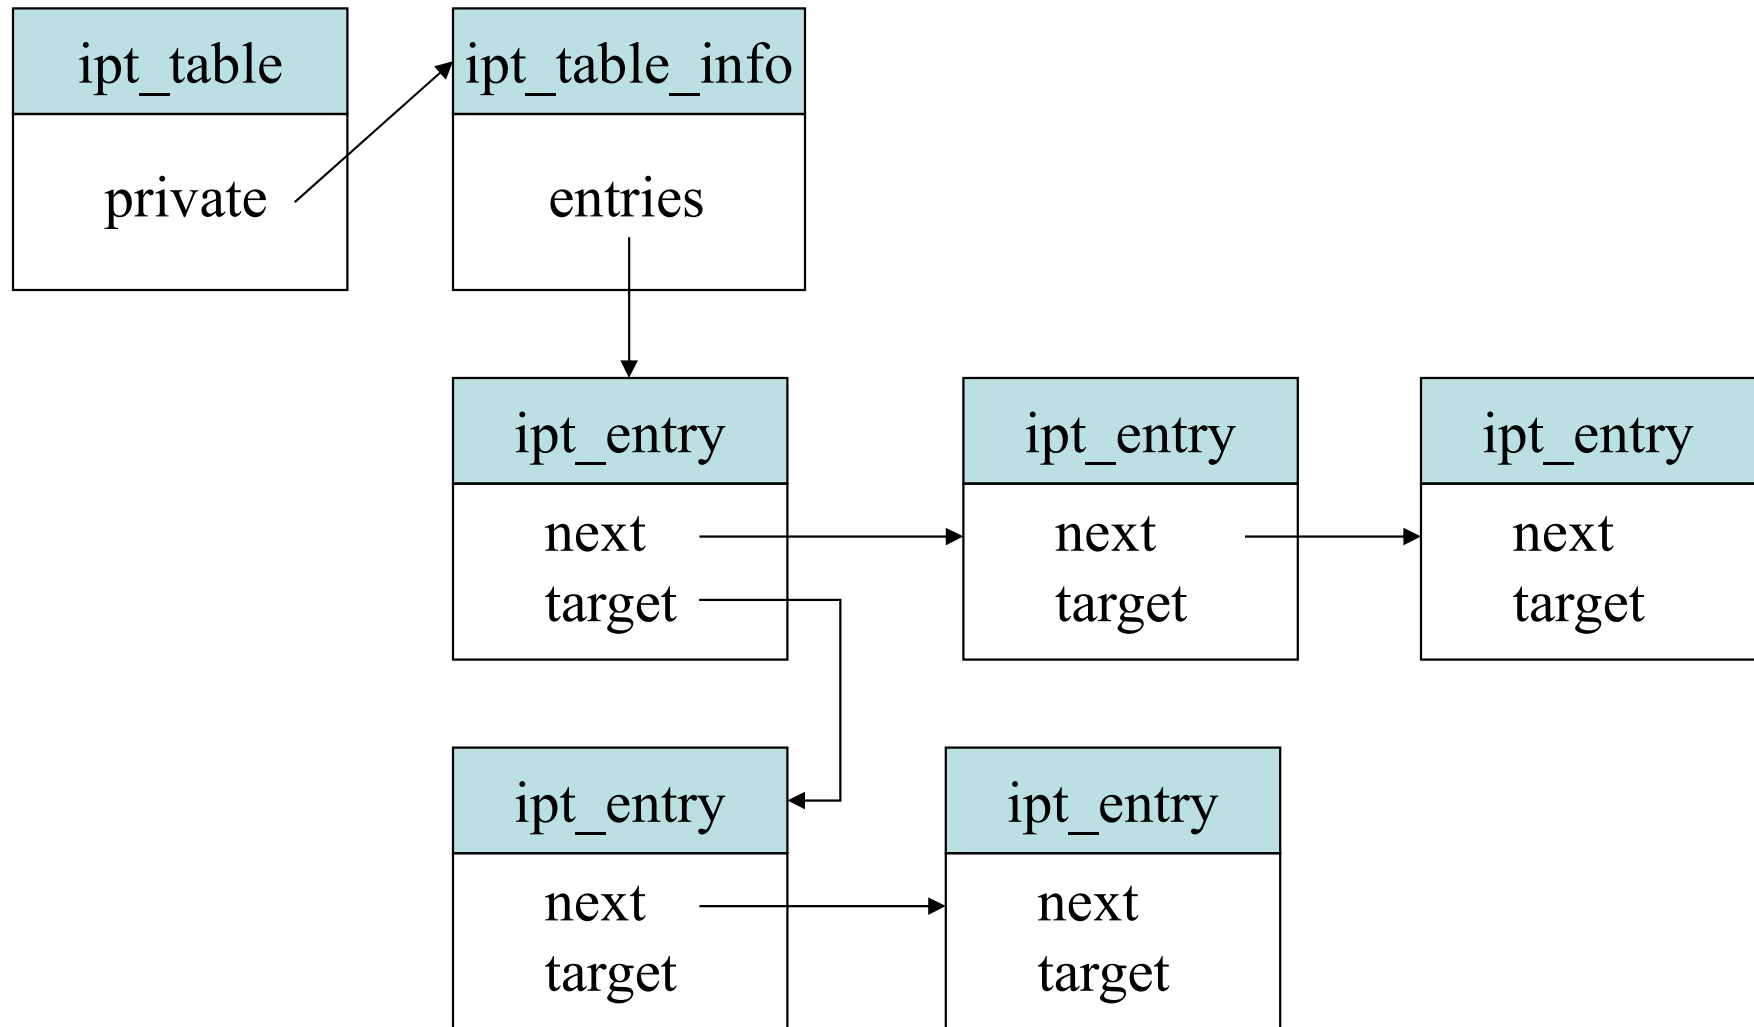
sk_buff



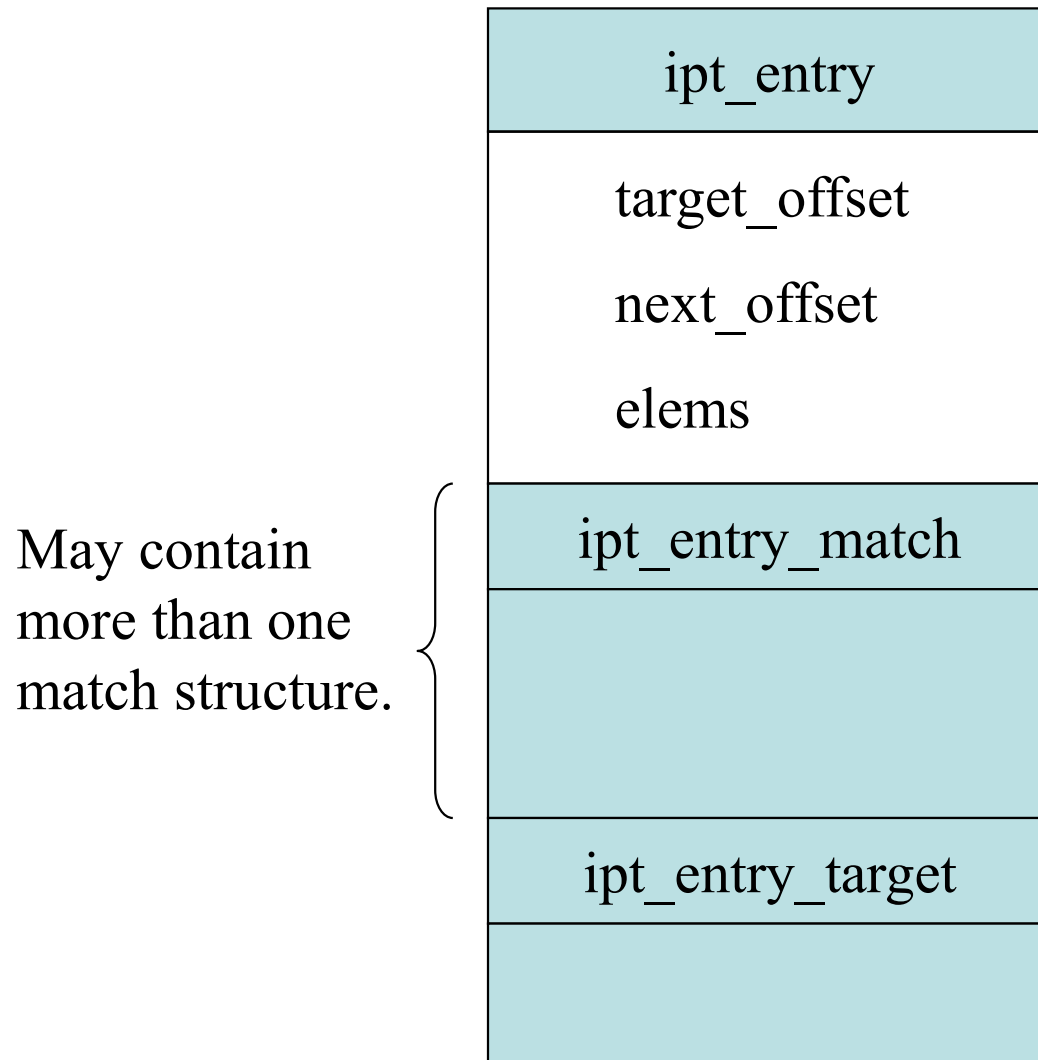
sk_buff

```
struct sk_buff {
    struct sk_buff  * next;           /* Next buffer in list */
    struct sk_buff  * prev;           /* Previous buffer in list */
    struct sock *sk;                   /* Socket we are owned by */
    struct timeval  stamp;              /* Time we arrived */
    struct net_device *dev;            /* I/O net device */
    /* Transport layer header */
    union
    {
        struct tcphdr    *th;
        struct udphdr    *uh;
        struct icmphdr  *icmph;
        struct iphdr    *ipiph;
    } h;
    /* Network layer header */
    union
    {
        struct iphdr     *iph;
        struct arphdr    *arph;
    } nh;
    ...
};
```

IP Tables Data Structures



struct ipt_entry



struct ipt_entry

```
struct ipt_entry
{
    /* Specifications for IP header we are to match */
    struct ipt_ip ip;
    /* Mark fields that rule examines. */
    unsigned int nfcache;

    /* Size of ipt_entry + matches */
    u_int16_t target_offset;
    /* Next ipt_entry: Size of ipt_entry + matches + target */
    u_int16_t next_offset;

    /* Back pointer */
    unsigned int comefrom;

    /* Packet and byte counters. */
    struct ipt_counters counters;

    /* The matches (if any), then the target. */
    unsigned char elems[0];
};
```

ipt_entry_match

struct ipt_entry_match contains

- Union of user and kernel structures.
 - Both contain match size.
 - Kernel part contains ptr to struct ipt_match.
- User-defined match information.

struct ipt_match

`name`: String that identifies this match.

`match`: Boolean function that determines whether the packet matched the rule or not.

`checkentry`: Boolean function that determines whether rule user attempted to enter was valid or not.

`destroy`: Called when rule deleted.

`me`: pointer to `THIS_MODULE`.

ipt_entry_target

struct ipt_entry_target contains

- Union of user and kernel structures.
 - Both contain target size.
 - Kernel part contains ptr to struct ipt_target.
- User-defined target information.

ipt_target

`name`: String that identifies target.

`target`: Returns a Netfilter action for the packet.

`checkentry`: Boolean function called when user attempts to enter this target.

`destroy`: Called when rule deleted.

`me`: pointer to `THIS_MODULE`.

Netfilter Actions

NF_ACCEPT

Allow packet to pass.

NF_DROP

Drop unacceptable packets

NF_STOLEN

Forget about packet.

NF_QUEUE

Queue packet for userspace program.

NF_REPEAT

Call this hook again.

Helper Functions

- `ipt_get_target()`
 - Returns pointer to target of a rule.
- `IPT_MATCH_ITERATE()`
 - Calls given fn for every match in given rule.
 - Function's 1st arg is struct `ipt_match_entry`.
 - Function returns zero for iteration to continue.
- `IPT_ALIGN()`
 - Calculates proper alignment of netfilter data structures.

ipt_do_table

Foreach `ipt_entry` in table:

 Foreach match in entry:

 If packet matches, call target.

 If target fn verdict is `IPT_RETURN`

 return to entry we were called from.

 Else if verdict `IPT_CONTINUE`

 continue on to next entry.

ipt_do_table

```
struct ipt_entry *e = get_entry(table_base,...)
unsigned int verdict = NF_DROP;
do {
    if (ip_packet_match(ip, indev, outdev, &e->ip, offset)) {
        struct ipt_entry_target *t;

        if (IPT_MATCH_ITERATE(e, do_match, ...) != 0)
            goto no_match;
        t = ipt_get_target(e); // then get verdict from target
    } else {
        no_match:
        e = (void *)e + e->next_offset;
    }
} while (!hotdrop);
if (hotdrop)
    return NF_DROP;
else return verdict;
```

do_match

```
int do_match(struct ipt_entry_match *m,
             const struct sk_buff *skb,
             const struct net_device *in,
             const struct net_device *out,
             int offset,
             const void *hdr,
             u_int16_t datalen,
             int *hotdrop)
{
    /* Stop iteration if it doesn't match */
    if (!m->u.kernel.match->match(skb, in, out, m->data,
                                   offset, hdr, datalen, hotdrop))
        return 1;
    else
        return 0;
}
```

Writing a Netfilter Extension

Userspace modifications: iptables

- Place libipt_foo.c in iptables/extensions.
- Add foo to iptables/extensions/Makefile.

Kernel modifications: netfilter

- Add ipt_foo.c in net/ipv4/netfilter.
- Add ipt_foo.h in include/linux/netfilter_ipv4.

Modifying iptables: libipt_foo.c

```
static struct iptables_match sctp = {
    .name          = "sctp",
    .version        = IPTABLES_VERSION,
    .size           = IPT_ALIGN(sizeof(struct
ipt_sctp_info)),
    .userspace_size = IPT_ALIGN(sizeof(struct
ipt_sctp_info)),
    .help           = &help,
    .init           = &init,
    .parse          = &parse,
    .final_check    = &final_check,
    .print          = &print,
    .save           = &save,
    .extra_opts     = opts
};

void _init(void) {    register_match(&sctp); }
```

Modifying iptables: libipt_foo.c

Functions in `iptables_match` handle options:

- help: `iptables -h`
- print: `iptables -L`
- parse: `iptables -[A|I|D|R]`
- save: `iptables-save`

Parsing sets `struct ipt_foo_info`

- Via `(*match)->data` field.
- Info struct defined in kernel include file.

Modifying Netfilter: ipt_foo.c

```
static struct ipt_match ipt_my_reg = {
    .list = { NULL, NULL }
    .name = "limit",
    .match = &match,
    .checkentry = &checkentry,
    .destroy = &destroy,
    .me = THIS_MODULE
};

static int __init limit_init(void)
{
    if (ipt_register_match(&ipt_my_reg))
        return -EINVAL;
    return 0;
}

static void __exit limit_fini(void)
{
    ipt_unregister_match(&ipt_my_reg);
}
```

References

1. Michael D. Bauer, *Linux Server Security*, 2nd edition, O'Reilly, 2005.
2. Christian Benvenuti, *Understanding Linux Network Internals*, O'Reilly, 2006.
3. Daniel P. Bovet and Marco Cesati, *Understanding the Linux Kernel*, 3rd edition, O'Reilly, 2005.
4. Thomas F. Herbert, *The Linux TCP/IP Stack*, Charles River Media, 2005.
5. Jennifer Hou, "Inside Netfilter,"
http://lion.cs.uiuc.edu/courses/cs498hou_spring05/lectures.html, 2005.
6. Lu-chuan Kung, "Netfilter Tutorial,"
http://lion.cs.uiuc.edu/courses/cs498hou_spring05/lectures.html, 2005.
7. Robert Love, *Linux Kernel Development*, 2nd edition, Prentice-Hall, 2005.
8. Rusty Russell, Linux World 2000 Netfilter Tutorial,
<http://www.netfilter.org/documentation/tutorials/lw-2000/tut.html>, 2000.
9. Rusty Russell and Harald Welte, Linux netfilter Hacking HOWTO,
<http://www.netfilter.org/documentation/HOWTO/netfilter-hacking-HOWTO.html>, 2002.
10. Rusty Russel, Linux Packet Filtering HOWTO,
<http://www.netfilter.org/documentation/HOWTO//packet-filtering-HOWTO.html>, 2002.